

Project Name : Prediction of Calories Burnt using XGBoost

Prediction of calories burnt Using Machine Learning

This project aims to predict the calories burnt during various physical activities using machine learning techniques. The dataset used for this project includes various features such as age, weight, duration of activity, etc.

Project Introduction

The project titled "Prediction of Calories Burnt using XGBoost" focuses on utilizing machine learning techniques to estimate the number of calories burnt during physical activities. Caloric expenditure is a key factor in weight management and overall health. Traditional methods of estimating calories burnt, such as using heart rate monitors or metabolic equations, can be cumbersome and often lack precision. This project aims to provide a more accurate and efficient method by leveraging the power of XGBoost, a scalable and high-performance gradient boosting framework.



PREDICT CALORIES BURNT

MACHINE LEARNING PROJECT



Problem Statement

Accurate prediction of calories burnt during physical activities is crucial for individuals who are keen on managing their weight, fitness, and overall health. Existing methods either require specialized equipment or depend on generalized equations that may not account for individual variability in metabolism and activity levels. The challenge is to develop a machine learning model that can predict calories burnt based on readily available data inputs, such as age, weight, height, gender, duration of activity, and type of activity

According to Jessica Simpson,"LAUGH A LOT. IT BURNS A LOT OF CALORIES"

- In this notebook, let's predict the calories burnt using XGBoost & let's have a healthy & a happier life.

DataSet Story

- User_ID : Unique identifier for each individual in the datasets
- Gender: Gender of the individual (e.g., 'Male' or 'Female').*
- Age: Age of the individual.
- Height: Height of the individual
- Weight: Weight of the individual.
- Duration: Duration of the physical activity or exercise session.
- Heart_Rate: Heart rate of the individual during the activity.
- Body_Temp: Body temperature of the individual.
- Calories: Calories burnt by an individual.
- y.*..**

IMPORTING LIBRARIES

```
import numpy as np import pandas as pd
import matplotlib.pyplot as plt import ydata_profiling as pp
import warnings
warnings.filterwarnings('ignore') import seaborn as sns
from sklearn.model_selection import train_test_split import xgboost
from xgboost import XGBRegressor from sklearn import metrics
```

LOADING THE DATA & PREPROCESSING THE DATA

```
calories_data = pd.read_csv("calories.csv") calories_data.head()
```

	User_ID	Calories
0	14733363	231.0
1	14861698	66.0
2	11179863	26.0
3	16180408	71.0

```
exercise_data = pd.read_csv("exercise.csv") exercise_data.head()
```

	User_ID	Gender	Age	Height	Weight	Duration	Heart_Rate	Body_Temp
--	---------	--------	-----	--------	--------	----------	------------	-----------

```

0    14733363   male  68  190.0  94.0    29.0   105.0
40.8
1    14861698   female 20  166.0  60.0    14.0    94.0
40.3
2    11179863   male  69  179.0  79.0    5.0     88.0
38.7
3    16180408   female 34  179.0  71.0    13.0   100.0
40.5
4    17771927   female 27  154.0  58.0    10.0    81.0
39.8

```

Inference:

- We can observe an Indirect relationship between heart rate & body temperature with respect to the exercise the individual does.

COMBINING THE EXERCISE DATA & CALORIES DATA

```
combined_data = pd.concat([exercise_data,calories_data['Calories']], axis=1) combined_data.head()
```

	User_ID	Gender	Age	Height	Weight	Duration	Heart_Rate
0	14733363	male	68	190.0	94.0	29.0	105.0
1	14861698	female	20	166.0	60.0	14.0	94.0
2	11179863	male	69	179.0	79.0	5.0	88.0
3	16180408	female	34	179.0	71.0	13.0	100.0
4	17771927	female	27	154.0	58.0	10.0	81.0

	Calories
0	231.0
1	66.0
2	26.0
3	71.0
4	35.0

#checking the number of rows & columns

```
combined_data.shape (15000, 9)
```

Hence our combined_data has 15,000 people with 9 different features

#getting the information about the data most importantly to know about any missing values

```
combined_data.info()
```

```
< class 'pandas.core.frame.DataFrame' >
RangeIndex: 15000 entries, 0 to 14999 Data columns (total 9
columns):
```

```

# Column    Non-Null Count Dtype 
--- 
0   User_ID    15000 non-null int64
1   Gender     15000 non-null object
2   Age        15000 non-null int64
3   Height     15000 non-null float64
4   Weight     15000 non-null float64
5   Duration   15000 non-null float64
6   Heart_Rate 15000 non-null float64
7   Body_Temp   15000 non-null float64 8   Calories   15000 non-null float64 dtypes: float64(6), int64(2), object(1) memory usage: 1.0+ MB

```

#checking the missing values combined_data.isnull().sum()

```

User_ID      0
Gender       0
Age         0
Height      0
Weight      0
Duration    0
Heart_Rate   0
Body_Temp   0   Calories   0
dtype: int64

```

ANALYSING THE DATA

#To get stastical measures about the data combined_data.describe()

```

User_ID      Age      Height     Weight
Duration \
count 1.500000e+04 15000.000000 15000.000000 15000.000000
15000.000000
mean 1.497736e+07 42.789800 174.465133 74.966867
15.530600
std 2.872851e+06 16.980264 14.258114 15.035657
8.319203
min 1.000116e+07 20.000000 123.000000 36.000000
1.000000
25% 1.247419e+07 28.000000 164.000000 63.000000
8.000000
50% 1.499728e+07 39.000000 175.000000 74.000000
16.000000
75% 1.744928e+07 56.000000 185.000000 87.000000
23.000000
max 1.999965e+07 79.000000 222.000000 132.000000 30.000000

Heart_Rate  Body_Temp   Calories count 15000.000000
15000.000000 15000.000000 mean 95.518533 40.025453
89.539533 std 9.583328 0.779230 62.456978 min 67.000000
37.100000 1.000000
25% 88.000000 39.600000 35.000000

```

```
50%    96.000000  40.200000  79.000000 75%    103.000000  
40.600000  138.000000 max    128.000000  41.500000 314.000000
```

Important inference


- Heart rate & body temperature is more when the person is doing exercise

```
pp.ProfileReport(combined_data)
```


```
{"model_id":"f837aa292b9d4795a178dc83a484fb02","version_major":2,"vers ion_minor":0}  
{"model_id":"cb048e6a2879499ca0c22fa667c8e550","version_major":2,"vers ion_minor":0}  
{"model_id":"5e2701fbfc7f42e293d2ef9f9a85d6d5","version_major":2,"vers ion_minor":0}  
< IPython.core.display.HTML object >
```

VISUALIZATION OF DATA

```
sns.set()  
  
%matplotlib inline  
sns.countplot(data=combined_data, x='Gender')  
  
# Show the plot plt.show()
```



From the above plot, we can observe that the data is evenly distributed(almost 7000) for both males and females.




```
#inorder to find the distribution of age column sns.distplot(combined_data['Age'])
```

From the above plot, we can observe that the Age column has the mean value of 42 and the data is positively skewed & most of the datapoints are in between 20 to 50.

```
#inorder to find the distribution of height column sns.distplot(combined_data['Height'])
```

```
<Axes: xlabel='Height', ylabel='Density'>
```




Hence,from the above plot we can observe that the height data is normally distributed with mean of 174cm.

```
#inorder to find the distribution of age column sns.distplot(combined_data['Age'])
```

```
<Axes: xlabel='Age', ylabel='Density'>
```


```
#inorder to find the distribution of weight column sns.distplot(combined_data['Weight'])
```

```
<Axes: xlabel='Weight', ylabel='Density'>
```




Hence, from the above plot we can observe that the weight data is also normally distributed with mean of 74 kg

```
sns.distplot(combined_data['Duration'])  
<Axes: xlabel='Duration', ylabel='Density'>
```




Hence, from the above plot we can observe that the duration data is a low peaked data with mean of 15mins

```
sns.distplot(combined_data['Heart_Rate'])  
<Axes: xlabel='Heart_Rate', ylabel='Density' >
```



Hence, from the above plot we can observe that the heart rate column is normally distributed with mean of 95 heart beats with several outliers.

```
sns.distplot(combined_data['Body_Temp'])  
<Axes: xlabel='Body_Temp', ylabel='Density' >
```




Hence, from the above plot we can observe that the temperature column is negatively skewed with mean of 40.


TO FIND THE CORRELATION IN THE DATA

Duration & calories are directly proportional=positively correlated


```
# Scatter plot for Age vs Calories
sns.scatterplot(x='Age', y='Calories', data=combined_data) plt.title('Scatter Plot: Age vs Calories') plt.show()
```




```
sns.boxplot(x='Gender',y='Heart_Rate',data=combined_data) plt.title("Boxplot: Gender vs Heart_Rate") plt.show()
```



```
# Scatter plot for Body Temp vs Heart Rate
sns.scatterplot(x='Body_Temp', y='Heart_Rate', data=combined_data) plt.title('Scatter Plot: Body Temp vs Heart Rate') plt.xlabel('Body Temperature') plt.ylabel('Heart Rate') plt.show()
```




```
# Bar plot for Duration vs Calories
sns.barplot(x='Duration',y='Calories',data=combined_data) plt.title("Bar graph
Calories vs Duration") plt.xlabel('Duration') plt.ylabel('Calories Burned') plt.show()
```




```
# Scatter plot for Age vs Weight
```


```
sns.scatterplot(x='Age', y='Weight', data=combined_data) plt.title('Scatter Plot: Age vs Weight') plt.show()
```



```
# Histogram for Duration  
sns.histplot(combined_data['Duration'], bins=30, kde=True, color='violet')  
plt.title('Histogram: Duration')  
plt.xlabel('Duration') plt.ylabel('Frequency')  
plt.show()
```



```
# Histogram for Calories Burned
sns.histplot(combined_data['Calories'], bins=30, kde=True, color='salmon')
plt.title('Histogram: Calories Burned') plt.xlabel('Calories Burned')
plt.ylabel('Frequency') plt.show()
```



```
#converting gender string data type to float
gender_mapping={'male':0,'female':1}

combined_data["Gender"] = combined_data["Gender"].map(gender_mapping)
combined_data
```

	User_ID	Gender	Age	Height	Weight	Duration	Heart_Rate	
Body_Temp \	14733363	0	68	190.0	94.0	29.0	105.0	
40.8	14861698	1	20	166.0	60.0	14.0	94.0	
40.3	11179863	0	69	179.0	79.0	5.0	88.0	
38.7	16180408	1	34	179.0	71.0	13.0	100.0	
40.5	17771927	1	27	154.0	58.0	10.0	81.0	
39.8	
...	14995	15644082	1	20	193.0	86.0	11.0	92.0
40.4	14996	17212577	1	27	165.0	65.0	6.0	85.0

39.2

```
14997    17271188     1   43  159.0  58.0   16.0   90.0
40.1
14998    18643037     0   78  193.0  97.0    2.0   84.0
38.3
14999    11751526     0   63  173.0  79.0   18.0   92.0
40.5
```

Calories

```
0           231.0
1           66.0
2           26.0
3           71.0
4           35.0 ... ...
14995      45.0
14996      23.0
14997      75.0
14998      11.0
14999      98.0
```


[15000 rows x 9 columns]

```
# df_numeric = combined_data[['Age', 'Height', 'Weight', 'Duration',
'Heart_Rate', 'Body_Temp']].copy()
# correlation_matrix = df_numeric.corr() correlation =
combined_data[['User_ID','Gender','Age','Height','Weight','Duration','Heart_Rate','Body_Temp']]
.corr()
```

Construction of heat map for the understanding of correlation

```
plt.figure(figsize=(10,10))
sns.heatmap(correlation, cbar=True, square=True, fmt='.1f', annot=True, annot_kws={'size':8},
cmap='Blues')

<Axes: >
```



CONVERSION OF TEXT DATA TO NUMERICAL VALUES

```
#combined_data.replace({'Gender':{'male':0,'female':1}},inplace=True) combined_data.head()

User_ID Gender Age Height Weight Duration Heart_Rate
Body_Temp \
0 14733363 0 68 190.0 94.0 29.0 105.0
40.8
1 14861698 1 20 166.0 60.0 14.0 94.0
40.3
2 11179863 0 69 179.0 79.0 5.0 88.0
38.7
3 16180408 1 34 179.0 71.0 13.0 100.0
40.5
4 17771927 1 27 154.0 58.0 10.0 81.0
39.8

Calories
0 231.0
1 66.0
2 26.0
3 71.0
4 35.0
```

SEPARATION OF FEATURES AND TARGET

```
X=combined_data.drop(['User_ID','Calories'],axis=1)
Y=combined_data['Calories'] print( X )
```

	Gender	Age	Height	Weight	Duration	Heart_Rate	Body_Temp
0	0	68	190.0	94.0	29.0	105.0	40.8
1	1	20	166.0	60.0	14.0	94.0	40.3
2	0	69	179.0	79.0	5.0	88.0	38.7
3	1	34	179.0	71.0	13.0	100.0	40.5
4	1	27	154.0	58.0	10.0	81.0	39.8
14995		1	20	193.0	86.0	11.0	92.0
14996		1	27	165.0	65.0	6.0	85.0
14997		1	43	159.0	58.0	16.0	90.0
14998		0	78	193.0	97.0	2.0	84.0
14999		0	63	173.0	79.0	18.0	92.0

[15000 rows x 7 columns] print(Y)

0	231.0
1	66.0
2	26.0
3	71.0
4	35.0
14995	45.0
14996	23.0
14997	75.0
14998	11.0

```
14999      98.0
Name: Calories, Length: 15000, dtype: float64
```

```
SPLIT THE DATA INTO TRAINING DATA & TEST DATA X_train,X_test,Y_train,Y_test =
train_test_split(X,Y,test_size=0.2,random_state=2) print( X.shape,X_train.shape,X_test.shape )
```

```
(15000, 7) (12000, 7) (3000, 7)
```

Hence, we can observe that

- X contains the original amount of data which is 15000
- X_train contains 80% of the data which is 12000
- X_test contains 20% of the data which is 3000

Regressor

```
#loading the model model =
XGBRegressor()

#training the model with X_train model.fit(X_train,Y_train)

XGBRegressor(base_score=0.5, booster='gbtree', callbacks=None,
             colsample_bylevel=1,
             colsample_bynode=1,
             colsample_bytree=1,
             early_stopping_rounds=None, enable_categorical=False,
             eval_metric=None, gamma=0, gpu_id=-1,
             grow_policy='depthwise',
             importance_type=None, interaction_constraints="",
             learning_rate=0.300000012,
             max_bin=256,
             max_cat_to_onehot=4,
             max_delta_step=0, max_depth=6, max_leaves=0,
             min_child_weight=1,
             missing=nan, monotone_constraints='()', n_estimators=100, n_jobs=0,
             num_parallel_tree=1, predictor='auto', random_state=0,
             reg_alpha=0, reg_lambda=1, ...)
```

EVALUATION

Evaluation is done by test data

PREDICTION OF THE DATA

```
calories_burnt_prediction = model.predict(X_test) print( calories_burnt_prediction )
```

```
[127.823784 226.00154 38.66253 ... 144.3636 22.767195
89.87375 ]
```

MEAN ABSOLUTE ERROR

```
MAE = metrics.mean_absolute_error(Y_test, calories_burnt_prediction) print("Mean Absolute Error
= ",MAE)
```

```
Mean Absolute Error = 1.4807048829992613
```

BUILDING A PREDICTIVE SYSTEM

- Building a predictive system inorder to find the calories burnt for the first individual from the dataset

```
input_data = (0,68,190.0,94.0,29.0,105.0,40.8,231)
```


```
print("The calories burnt for the first individual in the dataset is predicted as ",  
calories_burnt_prediction[0])
```

The calories burnt for the first individual in the dataset is predicted as 127.823784

```
print("Thus we have successfully predicted the calories burnt using XGBoost")
```

Thus we have successfully predicted the calories burnt using XGBoost
sns.regplot(x=Y_test,
y=calories_burnt_prediction)

```
plt.xlabel('Actual') plt.ylabel('Predicted')  
plt.show()
```



STAY SAFE

STAY HEALTHY